


CS/ECE 374 A ♦ Fall 2019
Midterm 2^y Problem 1 Solution

Clearly indicate the following structures in the directed graph G drawn below, or write NONE if the indicated structure does not exist.

1. A depth-first search tree rooted at a .

Solution: Here are three of *many* correct solutions:




Rubric: 2½ points. These are not the only correct solutions.

- No credit if the reported graph is not a spanning tree rooted at vertex a .
- -1 for each misplaced edge, compared to the closest correct solution

2. A breadth-first search tree rooted at a .

Solution: The correct solution is unique!




Rubric: 2½ points

- No credit if the reported graph is not a spanning tree rooted at vertex a .
- -1 for each misplaced edge

3. Circle each strong component.

Solution:




Rubric: 2½ points

- -1 for not circling the individual vertices e and h .
- No credit if circled subgraphs omit any vertex other than e or h .
- -1 for circling subgraphs ab and bcf , either instead of or in addition to $abcf$.
- No credit if any circled subgraph is not strongly connected.

4. Draw the strong-component graph of the example graph G .

Solution:



Rubric: 2½ points

- No credit if the reported graph is undirected or contains a directed cycle.
- -1 for each misplaced, omitted, or extra vertex or edge.
- No penalty for omitting the vertex labels.

CS/ECE 374 A ♦ Fall 2019
Midterm 2^y Problem 2 Solution

Eggsy has a map of the city in the form of an undirected graph G , whose vertices represent intersections and whose edges represent streets between them. A subset of the vertices are marked to indicate that the corresponding intersections are lit. Every edge e has a non-negative length $\ell(e)$. The graph has two special nodes s and t , which represent Eggsy's work and home, respectively. **Assume s and t are lit.**

Describe an algorithm that computes the shortest path in G from s to t that visits at most k unlit vertices, or correctly reports that no such path exists.

Solution: We construct a new *directed* graph $G' = (V', E')$ as follows:

- $V' = V \times \{0, 1, \dots, k\}$ — Each vertex (v, i) represents Eggsy reaching vertex v after visiting exactly i unlit vertices (including v , if v is unlit).
- E' contains two types of directed edges:
 - Edges into lit vertices: $\{(u, i) \rightarrow (v, i) \mid uv \in E \text{ and } v \text{ is lit}\}$
 - Edges into unlit vertices: $\{(u, i) \rightarrow (v, i + 1) \mid uv \in E \text{ and } v \text{ is unlit}\}$
- Each edge $(u, i) \rightarrow (v, j)$ has weight $\ell(uv)$.

We need to compute the minimum distance from $(s, 0)$ to any vertex of the form (t, i) . We can compute all such distances by running Dijkstra's algorithm once from $(s, 0)$, after which we can find the minimum distance in $O(k)$ additional time. The resulting algorithm runs in $O(E' \log V') = O(kE \log(kV)) = O(kE \log V)$ time. ■

Rubric: 10 points: standard graph reduction rubric. −1 for vanishing k into the $O()$ notation.

CS/ECE 374 A ♦ Fall 2019
Midterm 2^y Problem 3 Solution

An undirected graph $G = (V, E)$ is **bipartite** if its vertices can be partitioned into two subsets L and R , such that every edge in E has one endpoint in L and one endpoint in R . Describe and analyze an algorithm to determine, given an undirected graph G as input, whether G is bipartite. [Hint: Every tree is bipartite.]

Solution: Let G be the input graph. The algorithm runs in three phases:

- First, compute a spanning tree T of G using whatever-first search, starting at an arbitrary vertex s . (If G is disconnected, consider each component of G separately, and return TRUE if and only if every component is bipartite.)
- Color s black. Using a preorder traversal of T , color every vertex except s the opposite color of its parent.
- Finally, scan through all edges in G by brute force; if any edge joins two vertices with the same color, return FALSE, and if all edges join vertices of both colors, return TRUE.

This algorithm runs in $O(V + E)$ time. ■

Solution: Let G be the input graph. The algorithm runs in two phases:

- Perform a breadth-first search of the graph, starting at an arbitrary source vertex s , labeling every vertex v with its shortest-path distance $\text{dist}(v)$ from s . (If G is not connected, consider each component of G separately, and return TRUE if and only if every component is bipartite.)
- Then for every edge uv in the graph, if $\text{dist}(u) = \text{dist}(v)$, return FALSE; if every edge joins vertices with two different distances, return TRUE.

This algorithm runs in $O(V + E)$ time.

This algorithm relies on a specific property of shortest-path distances in *unweighted, undirected* graphs: For every source vertex s and every edge uv , the distance to from s to u and the distance from s to v differ by at most 1. ■

Solution: The following variant of depth-first search either colors every vertex black or white, so that every edge has one endpoint of each color, or fails because the input graph is not bipartite.

```

2COLOR(G):
  for all vertices v
    v.color ← NONE
  for all vertices v
    if v.color = NONE
      2COLORDFS(v, WHITE)

```

```

2COLORDFS(v, mycolor):
  if mycolor = WHITE
    yourcolor ← BLACK
  else
    yourcolor ← WHITE
  v.color ← mycolor
  for each edge vw
    if w.color = mycolor
      fail gracefully
    else if w.color = NONE
      2COLORDFS(w, yourcolor)

```

The algorithm runs in $O(V + E)$ time. ■

Solution: Assume without loss of generality that the input graph $G = (V, E)$ is connected. (Otherwise, consider each component of G separately, and return TRUE if and only if every component of G is bipartite.) Construct a new undirected graph $G' = (V', E')$ as follows:

- $V' = V \times \{0, 1\}$
- $E' = \{(u, 0)(v, 1) \mid uv \in E\} \cup \{(u, 1)(v, 0) \mid uv \in E\}$

G' has exactly twice as many vertices and twice as many edges as G . Now perform a whatever-first search in G' starting at an arbitrary vertex. If this search marks every vertex of G' , then report that G is *not* bipartite; if at least one vertex remains unmarked, report that G is bipartite. This algorithm runs in $O(V + E)$ time.

Suppose G is not bipartite. Then G contains an odd cycle $v_1 \rightarrow v_2 \rightarrow \dots \rightarrow v_{2k+1} \rightarrow v_1$, and therefore H contains the cycle

$$(v_1, 0) \rightarrow (v_2, 1) \rightarrow \dots \rightarrow (v_{2k+1}, 0) \rightarrow (v_1, 1) \rightarrow (v_2, 0) \rightarrow \dots \rightarrow (v_{2k+1}, 1) \rightarrow (v_1, 0).$$

Because G is connected, G contains a path from v_1 to every other vertex u of G . It follows that H contains a path from $(v_1, 0)$ to every vertex (u, i) of H , obtained either by following some $v_1 \rightsquigarrow u$ path in G directly, or by going around the cycle to $(v_1, 1)$ and then following some $v_1 \rightsquigarrow u$ path in G . We conclude that H is connected.

On the other hand, suppose G is bipartite. Color the vertices of G alternately black and white. Consider any two vertices u and v of G . Any path between $(u, 0)$ and $(v, 0)$ must have even length; but any path between a black vertex and a white vertex must have odd length. Thus, $(u, 0)$ and $(v, 0)$ lie in the same component of H if and only if u and v have the same color. Similarly, $(u, 0)$ and $(v, 1)$ lie in the same component of H if and only if u and v have different colors. We conclude that H consists of two disjoint copies of G , one with $0 = \text{white}$ and $1 = \text{black}$, the other with $0 = \text{black}$ and $1 = \text{white}$. ■

Rubric: 10 points = 8 points for the algorithm (if recursive: 2 point for base case + 6 points for recursive case) + 2 point for time analysis. These are not the only correct solutions. No penalty if the algorithm works only for connected graphs. A proof of correctness is not required. Max 7 points for an algorithm that runs in $O(E \log V)$ time; max 4 points for any slower correct algorithm.

CS/ECE 374 A ♦ Fall 2019
Midterm 2^y Problem 4 Solution

Describe an efficient algorithm that computes the largest number of students that Satya can host for testing in a sequence of n rooms without using three consecutive rooms. The input to your algorithm is an array $S[1..n]$, where each $S[i]$ is the number of students that can fit in room i . (See the question handout for more details.)

Solution: For any index i and any integer $c \in \{0, 1, 2\}$, let $MaxStudents(i, c)$ denote the maximum number of students that can be hosted for testing in rooms i through n , assuming exactly c rooms immediately before room i are already being used. We need to compute $MaxStudents(1, 0)$.

This function can be described by the following recurrence:

$$MaxStudents(i, c) = \begin{cases} 0 & \text{if } i > n \\ MaxStudents(i+1, 0) & \text{if } c = 2 \\ \max \begin{cases} S[i] + MaxStudents(i+1, c+1) \\ MaxStudents(i+1, 0) \end{cases} & \text{otherwise} \end{cases}$$

We can memoize this function into an array $MaxStudents[1..n, 0..2]$, which we can fill by decreasing i in the outer loop and considering c in any order in the inner loop. The resulting algorithm runs in $O(n)$ time. ■

Solution: For any index i , let $MaxStudents(i)$ denote the maximum number of students that can be hosted for testing in rooms i through n , assuming room $i-1$ is *not* occupied. We need to compute $MaxStudents(1)$.

This function can be described by the following recurrence:

$$MaxStudents(i) = \begin{cases} 0 & \text{if } i > n \\ S[n] & \text{if } i = n \\ S[n-1] + S[n] & \text{if } i = n-1 \\ \max \begin{cases} MaxStudents(i+1) \\ S[i] + MaxStudents(i+2) \\ S[i] + S[i+1] + MaxStudents(i+3) \end{cases} & \text{otherwise} \end{cases}$$

We can memoize this function into an array $MaxStudents[1..n]$, which we can fill by decreasing i in $O(n)$ time. ■

Rubric: 10 points: standard dynamic programming rubric. These are not the only correct solutions.

CS/ECE 374 A ♦ Fall 2019
Midterm 2^y Problem 5 Solution

Describe and analyze an algorithm to compute the number of maximal points in a given point set P in $O(n \log n)$ time. The input to your algorithm is a pair of arrays $X[1..n]$ and $Y[1..n]$ containing the x - and y -coordinates of the points in P . (See the question handout for more details and an example.)

All model solutions assume the given points have distinct x - and y -coordinates.

Solution (sort and scan backward): The following algorithm runs in $O(n \log n)$ time; the running time is dominated by the initial sort.

NUMMAXIMAL(X, Y) :
Sort X and permute Y to match *⟨⟨= Sort the points from left to right⟩⟩*
 ⟨⟨Scan right to left, counting new maximum y -coordinates⟩⟩
 $count \leftarrow 0$
 $maxy \leftarrow -\infty$
for $i \leftarrow n$ down to 1
 If $Y[i] > maxy$
 $count \leftarrow count + 1$
 $maxy \leftarrow Y[i]$
return $count$

■

Solution (sort and scan forward): The following algorithm runs in $O(n \log n)$ time; the running time is dominated by the initial sort.

NUMMAXIMAL(X, Y) :
Sort X and permute Y to match *⟨⟨= Sort the points from left to right⟩⟩*
 ⟨⟨Scan left to right⟩⟩
 $S \leftarrow$ new stack *⟨⟨maintain staircase of first i points⟩⟩*
PUSH(S, ∞) *⟨⟨sentinel⟩⟩*
for $i \leftarrow 1$ down to ∞
 while $Y[i] > \text{TOP}(S)$
 POP(S)
 PUSH($S, Y[i]$)
return SIZE(S) - 1

■

Solution (like quicksort): To simplify notation, assume the points are actually given in a single array $P[1..n]$, where each $P[i]$ stores a point with x -coordinate $P[i].x$ and y -coordinate $P[i].y$. We can convert from two separate coordinate arrays to this array of records in $O(n)$ time by brute force.

The following divide-and-conquer algorithm actually computes an array containing all maximal points in P , sorted by increasing x -coordinate, in **$O(n \log n)$ time**. The running time is dominated by the initial sort. We can compute the *number* of points in this list in $O(n)$ additional time.

STAIRCASE($P[1..n]$):
 Sort P by increasing x -coordinate
 QUICKSTAIR($P[1..n]$)

QUICKSTAIR($P[1..n]$):
 if $n = 1$
 return $P[1]$
« Recursively construct staircases »
« of the left and right halves of P »
 $L[1..l] \leftarrow \text{QUICKSTAIR}(P[1..\lceil n/2 \rceil])$
 $R[1..r] \leftarrow \text{QUICKSTAIR}(P[\lfloor n/2 \rfloor + 1..n])$
« Discard points on the left staircase that lie »
« below the first point on the right staircase »
 while $L[l].y \leq R[1].y$
 $l \leftarrow l - 1$
« Copy to an output array and return »
 for $i \leftarrow 1$ to l
 $M[i] \leftarrow L[i]$
 for $i \leftarrow 1$ to r
 $M[l + i] \leftarrow R[i]$
 return $M[1..l + r]$

The running time of QUICKSTAIR obeys the standard *mergesort* recurrence $T(n) = 2T(n/2) + O(n)$, so QUICKSTAIR runs in $O(n \log n)$ time. ■

Solution (like mergesort): To simplify notation, assume the points are actually given in a single array $P[1..n]$, where each $P[i]$ stores a point with x -coordinate $P[i].x$ and y -coordinate $P[i].y$. We can convert from two separate coordinate arrays to this array of records in $O(n)$ time by brute force.

The following algorithm actually computes an array containing all maximal points in P , sorted from left to right. Unlike the previous solutions, this algorithm does *not* begin by sorting the points. We can compute the *number* of points in this list in $O(n)$ additional time.

```

MERGESTAIR( $P[1..n]$ ):
  if  $n = 1$ 
    return  $P[1]$ 

  « Recursively construct staircases »
  « of two arbitrary halves of  $P$  »
   $A[1..a] \leftarrow \text{MERGESTAIR}(P[1..\lceil n/2 \rceil])$ 
   $B[1..b] \leftarrow \text{MERGESTAIR}(P[\lceil n/2 \rceil + 1..n])$ 

  « Merge the overlapping staircases from left to right »
   $m \leftarrow 0$  « number of output points »
   $i \leftarrow 1; j \leftarrow 1$ 
  while  $i + j < a + b$ 
    if  $j = b$  « we've finished scanning B; copy next point from A »
       $m \leftarrow m + 1; M[m] \leftarrow A[i]$ 
       $i \leftarrow i + 1$ 
    else if  $i = a$  « we've finished scanning A; copy next point from B »
       $m \leftarrow m + 1; M[m] \leftarrow B[j]$ 
       $j \leftarrow j + 1$ 
    else if  $A[i].x < B[j].x$  « next unscanned point is in A »
      if  $A[i].y > B[j].y$  « and not hidden by next point in B »
         $m \leftarrow m + 1; M[m] \leftarrow A[i]$ 
         $i \leftarrow i + 1$ 
      else «  $A[i].x > B[j].x$  » « next unscanned point is in B »
        if  $B[j].y > A[i].y$  « and not hidden by next point in A »
           $m \leftarrow m + 1; M[m] \leftarrow B[j]$ 
           $j \leftarrow j + 1$ 
  return  $M[1..m]$ 

```

The running time of MERGESTAIR obeys the standard mergesort recurrence $T(n) = 2T(n/2) + O(n)$, so MERGESTAIR runs in $O(n \log n)$ time. ■

Rubric: 10 points = 7 points for algorithm + 3 points for time bound. These are not the only correct solutions. Max 5 points for an $O(n^2)$ -time algorithm; max 3 points for anything slower. No penalty for misbehavior when x - or y -coordinates are not distinct. These are not the only correct solutions.