

1. Suppose that you have just finished computing the array  $dist[1..V, 1..V]$  of shortest-path distances between **all** pairs of vertices in an edge-weighted directed graph  $G$ . Unfortunately, you discover that you incorrectly entered the weight of a single edge  $u \rightarrow v$ , so all that precious CPU time was wasted. Or was it? Maybe your distances are correct after all!

In each of the following problems, let  $w(u \rightarrow v)$  denote the weight that you used in your distance computation, and let  $w'(u \rightarrow v)$  denote the correct weight of  $u \rightarrow v$ .

- (a) Suppose  $w(u \rightarrow v) > w'(u \rightarrow v)$ ; that is, the weight you used for  $u \rightarrow v$  was *larger* than its true weight. Describe an algorithm that repairs the distance array in  $O(V^2)$  time under this assumption. [Hint: For every pair of vertices  $x$  and  $y$ , either  $u \rightarrow v$  is on the shortest path from  $x$  to  $y$  or it isn't.]

**Solution:** Consider any two vertices  $s$  and  $t$ . If the true shortest path from  $s$  to  $t$  contains the mistake edge  $u \rightarrow v$ , then its length is  $dist[s, u] + w'(u \rightarrow v) + dist[v, t]$ . If the true shortest path from  $s$  to  $t$  does not contain the mistaken edge  $u \rightarrow v$ , then  $dist[s, t]$  is correct.

```

REPAIRDISTANCES( $dist, w'(u \rightarrow v)$ ):
  for every vertex  $s$ 
    for every vertex  $t$ 
      if  $dist[s, t] > dist[s, u] + w'(u \rightarrow v) + dist[v, t]$ 
         $dist[s, t] \leftarrow dist[s, u] + w'(u \rightarrow v) + dist[v, t]$ 

```

■

- (b) Maybe even that was too much work. Describe an algorithm that determines whether your original distance array is actually correct in  $O(1)$  time, again assuming that  $w(u \rightarrow v) > w'(u \rightarrow v)$ . [Hint: Either  $u \rightarrow v$  is the shortest path from  $u$  to  $v$  or it isn't.]

**Solution:** The edge  $u \rightarrow v$  appears in any shortest path if and only if  $u \rightarrow v$  itself is a shortest path from  $u$  to  $v$ . Thus, if  $u \rightarrow v$  is *not* the unique shortest path from  $u$  to  $v$  after fixing its weight, then all shortest paths can avoid  $u \rightarrow v$ , which means all the old distances are correct. On the other hand, if  $dist[u, v] > w'(u \rightarrow v)$ , then at least  $dist[u, v]$  is incorrect.

```

CHECKDISTANCES( $dist, w'(u \rightarrow v)$ ):
  if  $dist[u, v] \leq w'(u \rightarrow v)$ 
    return TRUE
  else
    return FALSE

```

■

- (c) **To think about later:** Describe an algorithm that determines in  $O(VE)$  time whether your distance array is actually correct, even if  $w(u \rightarrow v) < w'(u \rightarrow v)$ .

**Solution:** If  $w(u \rightarrow v) < w'(u \rightarrow v)$ , we need to compute the correct shortest-path distance from  $u$  to  $v$ . If this new distance is equal to the old value  $dist[u, v]$ , then  $u \rightarrow v$  was not a shortest path under the old weights, so all (old and new) shortest paths avoid  $u \rightarrow v$ , so all the old distances are correct. Otherwise, at least the distance  $dist[u, v]$  is incorrect.

```

CHECKDISTANCES( $dist, w'(u \rightarrow v)$ ):
    compute  $dist'[u, v]$  via Bellman-Ford
    if  $dist'[u, v] = dist[u, v]$ 
        return TRUE
    else
        return FALSE

```

■

- (d) **To think about later:** Argue that when  $w(u \rightarrow v) < w'(u \rightarrow v)$ , repairing the distance array *requires* recomputing shortest paths from scratch, at least in the worst case.

**Solution:** Let  $G$  be an arbitrary edge-weighted directed graph. Construct a new graph from  $H$  by adding two vertices  $u$  and  $v$ , edges  $x \rightarrow u$  and  $v \rightarrow x$  with length 0 for every vertex  $x$  in  $G$ , and an edge  $u \rightarrow v$  with weight  $-\infty$ . Then every shortest path in  $H$  has length  $-\infty$ , because it contains the edge  $u \rightarrow v$  and at most two other edges  $x \rightarrow u$  and  $v \rightarrow y$ . In particular, the lengths of the edges in  $G$  are utterly irrelevant.

Now if we set  $w'(u \rightarrow v) = \infty$ , then the new shortest path in  $H$  between two nodes of  $G$  is just their shortest path in  $G$ . But we have absolutely no information about shortest paths in  $G$ ; all we have is a distance array full of incorrect  $-\infty$ s! We have no choice but to recompute all shortest paths in  $G$  from scratch. ■

2. Suppose  $n$  different currencies are traded in your currency market. You are given the matrix  $R[1..n]$  of exchange rates between every pair of currencies; for each  $i$  and  $j$ , one unit of currency  $i$  can be traded for  $R[i, j]$  units of currency  $j$ . (Do *not* assume that  $R[i, j] \cdot R[j, i] = 1$ .)
- (a) Describe an algorithm that returns an array  $V[1..n]$ , where  $V[i]$  is the maximum amount of currency  $i$  that you can obtain by trading, starting with one unit of currency 1, assuming there are no arbitrage cycles.

**Solution:** Construct a complete graph  $G$  on  $n$  vertices with edge weights

$$w(i \rightarrow j) := -\lg R[i, j].$$

Any sequence of trades that starts with one unit of currency  $i$  and ends with  $M$  units of currency  $j$  corresponds to a path in  $G$  from vertex  $i$  to vertex  $j$  with length  $-\lg M$ . Conversely, any path of length  $\ell$  from vertex  $i$  to vertex  $j$  corresponds to a sequence of trades that starts with one unit of currency  $i$  and ends with  $2^{-\ell}$  units of currency  $j$ . In particular, a negative cycle in  $G$  would correspond to an arbitrage cycle; thus,  $G$  has no negative cycles.

Compute the shortest paths from vertex 1 to every other vertex in  $G$ , using Bellman-Ford, because some edge weights may be negative. Bellman-Ford runs in  $O(VE) = O(n^3)$  time. Finally, for each  $j$ , let  $V[j] = 2^{-\text{dist}(j)}$ , where  $\text{dist}(j)$  is the shortest-path distance from vertex 1 to vertex  $j$  computed by Bellman-Ford. Computing the output array  $V[1..n]$  requires only  $O(n)$  additional time.

Alternatively, if we don't like logs and exponents, we can modify Bellman-Ford to *multiply* edge lengths instead of adding them, and to reverse the direction of all comparisons. Here is the resulting algorithm, which clearly runs in  $O(n^3)$  time:

```

BELLMANFORDTRADING( $R[1..n, 1..n]$ )
   $V[1] \leftarrow 1$ 
  for  $i \leftarrow 2$  to  $n$ 
     $V[i] \leftarrow 0$ 
  for  $k \leftarrow 1$  to  $n-1$ 
    for  $i \leftarrow 1$  to  $n$ 
      for  $j \leftarrow 1$  to  $n$ 
        if  $V[j] \leq V[i] \cdot R[i, j]$ 
           $V[j] \leftarrow V[i] \cdot R[i, j]$ 
  return  $V[1..n]$ 

```

[I am assuming here that each arithmetic operation takes only  $O(1)$  time.] ■

- (b) Describe an algorithm to determine whether the given array of currency exchange rates creates an arbitrage cycle.

**Solution:** One more iteration of Bellman-Ford detects negative cycles, so we can use almost the same algorithm as in part (a).

```

BELLMANFORDARBITRAGE( $R[1..n, 1..n]$ )
   $V[1] \leftarrow 1$ 
  for  $i \leftarrow 2$  to  $n$ 
     $V[i] \leftarrow 0$ 
  for  $k \leftarrow 1$  to  $n - 1$ 
    for  $i \leftarrow 1$  to  $n$ 
      for  $j \leftarrow 1$  to  $n$ 
        if  $V[j] \leq V[i] \cdot R[i, j]$ 
           $V[j] \leftarrow V[i] \cdot R[i, j]$ 
  for  $i \leftarrow 1$  to  $n$ 
    for  $j \leftarrow 1$  to  $n$ 
      if  $V[j] \leq V[i] \cdot R[i, j]$ 
        return TRUE
  return FALSE

```

■

- \* (c) **To think about later:** Modify your algorithm from part (b) to actually return an arbitrage cycle, if such a cycle exists.

**Solution:** We further modify Bellman-Ford to maintain predecessor edges, exactly as described in the lecture notes. Then if there is a negative cycle in the graph, at least one such cycle is described by the predecessor edges; conversely, if the predecessor edges induce a cycle, the total weight of that cycle must be negative. Thus, we can find a negative weight cycle in  $O(V + E) = O(n^2)$  additional time using an obvious modification the ISACYCLIC algorithm in the notes.

Of course, the two underlying claims require proof.

**Claim 1.** *If there is a negative cycle in the graph, then after  $n$  iterations of Bellman-Ford, there is a cycle in the graph of predecessor edges.*

**Proof:** To simplify discussion, assume that every other vertex is reachable from  $s$  (as in the case in our arbitrage problems). If there is no negative cycles in the graph, then every vertex *except  $s$  itself* has an incoming predecessor edge when Bellman-Ford halts. If there is a negative cycle containing  $s$ , the algorithm will relax some edge *into*  $s$ , giving  $s$  an incoming predecessor edge. At that point, *every* vertex has an incoming predecessor edge. Thus, if we walking backward along those edges we will never get stuck; we must eventually repeat a vertex.

More generally, let  $N$  be the set of vertices reachable from  $s$  that lie on a negative cycle; obviously every vertex in such a negative cycle must lie in  $N$ . After  $n$  iterations of Bellman-Ford, the predecessor of each vertex in  $N$  is also in  $N$ . It follows that there must be a cycle among the predecessor edges in  $N$ . □

**Claim 2.** *If there is a cycle in the graph of predecessor edges after Bellman-Ford halts, the total weight of that cycle is negative.*

**Proof:** Consider a predecessor cycle  $C = v_0 \rightarrow v_1 \rightarrow \dots \rightarrow v_{\ell-1} \rightarrow v_0$ , where for each index  $i$ , we have  $\text{pred}(v_i) = (v_{i-1} \bmod \ell)$ . (I'll omit the “mod  $\ell$ ” from now on.) For each index  $i$ , define

$$\tilde{w}(v_{i-1} \rightarrow v_i) = w(v_{i-1} \rightarrow v_i) - \text{dist}(v_i) + \text{dist}(v_{i-1})$$

Just after the last time  $v_{i-1} \rightarrow v_i$  was relaxed, we had  $\tilde{w}(v_{i-1} \rightarrow v_i) = 0$ ; since that time,  $\text{dist}(v_i)$  has not changed and  $\text{dist}(v_{i-1})$  has not increased. It follows that  $\tilde{w}(v_{i-1} \rightarrow v_i) \leq 0$  for all  $i$ .

Suppose  $v_{i-1} \rightarrow v_i$  was the last edge in  $C$  to be relaxed. That relaxation decreased  $\text{dist}(v_i)$ , and therefore decreased  $\tilde{w}(v_i \rightarrow v_{i+1})$ , so we must have  $\tilde{w}(v_{i-1} \rightarrow v_i) < 0$ . (Equivalently,  $v_i \rightarrow v_{i+1}$  must be tense!)

Finally, we can express the total length of  $C$  in terms of the adjusted weights  $\tilde{w}$  as follows:

$$\begin{aligned} \sum_{i=0}^{\ell-1} w(v_{i-1} \rightarrow v_i) &= \sum_{i=0}^{\ell-1} (\tilde{w}(v_{i-1} \rightarrow v_i) + \text{dist}(v_i) - \text{dist}(v_{i-1})) \\ &= \sum_{i=0}^{\ell-1} \tilde{w}(v_{i-1} \rightarrow v_i) + \sum_{i=0}^{\ell-1} \text{dist}(v_i) - \sum_{i=0}^{\ell-1} \text{dist}(v_{i-1}) \\ &= \sum_{i=0}^{\ell-1} \tilde{w}(v_{i-1} \rightarrow v_i) \end{aligned}$$

Every term in the final sum is non-positive, and at least one term is negative. We conclude that  $C$  is a negative cycle. □

■